

# Introduction to Machine Learning (SS 2024)

## Project: Programming Project

### Author 1

Last name: Rieser  
First name: David  
Matrikel Nr.: 12141689

### Author 2

Last name: Sillaber  
First name: Cedric  
Matrikel Nr.: 12211124

## I. INTRODUCTION

The task of this project is to implement machine learning models to detect fraudulent transactions. The problem type is a binary classification task. Based on the timestamp, amount and 28 additional features the model should predict if a transaction is fraudulent or not. The dataset consists of 227845 instances with 30 features and flag indicating if the transaction is fraudulent or not. The dataset is highly imbalanced with 0.173% of the transactions being fraudulent. There is no information about the 28 features, so it is not clear what they represent.

As there the dataset is labelled, we can use supervised learning to solve this problem. We decided to use three methods for this problem. The first choice was to use a neural network for this task. Neural networks can extract the most important features from the data and are able to learn complex patterns. The second choice was a decision tree. In order to achieve even better results, we also implemented a random forest classification.

## II. IMPLEMENTATION / ML PROCESS

As mentioned in the introduction the dataset is highly imbalanced. There are 227845 datapoints and less than 500 are fraudulent. This is about 0.173% of the dataset. A naive implementation of the neural network was not able to detect frauds correctly. With an accuracy of 99.83% most likely the frauds were not detected. In order to solve this problem, we decided to do simple data augmentation. As other methods were not adequate, we simply duplicated the fraudulent transactions. This was done the following way: The dataset was split into train and validation set. Then the fraudulent transactions in the train set were duplicated randomly by a factor. Using this augmentation method it was possible to give the frauds more weight in the train set, without changing the validation set. Additionally we scaled the data using a normalization method. This functionality is achieved by calling scikit-learn's 'StandardScaler'.

For the problem we used two Methods: Multilayer Perceptron and Decision Tree.

### A. Multilayer Perceptron

The Neural Network is based on 'pytorch's 'nn.Module'. For the Network there are three layers. A input layer with 30 nodes, a hidden layer with dimensionality 94 and a learning rate of 0.0004301150216793739. The activation function is a ReLU function. The output layer has one node and a sigmoid activation function.

A Neural Network is suitable for this problem as it can learn complex patterns and extract the most important features from the data. We don't have any information about the features, so a Neural Network is a good choice. The programmer does not have to make any choices concerning the features. The dataset contains labeled data, so supervised learning and thus neural networks seem a good option.

The choice of hyperparameters was really important and made a difference. In this case the hyperparameters where the factor used for the data augmentation, the number of nodes in the hidden layer, learning rate and learning epochs. To find the best hyperparameters we used brute force method to train using randomly selected parameters. The parameters were randomly selected from a specified range. Given the results the best hyperparameters were chosen. The final hyperparameters are: factor for data augmentation: 50, number of nodes in hidden layer: 94, learning rate: 0.0004301150216793739, learning epochs: 36.

### B. Decision Tree

The Decision Tree is based on the 'sklearn' library.

Decision Trees are suitable for this problem as it can easily divide the multi-dimensional space given by the amount of features into distinct sections each with a given probability of being a fraud. For this the decision

Unlike the neural network the structure of the decision tree is determined from the learning process and thus the choice of hyperparameters is by far the most important part. After some testing with different hyperparameters we came to the following conclusion: Most of the hyperparameters result in the decision tree predicting only non-frauds. This happens because these parameters would usually expect a balanced dataset with multiple frauds per leaf and cannot work correctly with the miniscule amount of frauds.

This left us with the following hyperparameters: The maximum depth of the tree, the criterion on when to split the nodes, how to split the nodes, and the maximum amount of features.

### C. Random Forest

Further we tested a modification of decision trees - the random forest classification. This method is based on the decision tree, but uses multiple trees to improve the accuracy. The accuracy was about 99.959% and the model was able to detect the frauds. This is about the same as the accuracy of the decision tree.

## III. RESULTS

For the MLP model the final accuracy on the validation set was 99.9414%, while the train accuracy is only about 98%. This rather bad accuracy on the training set is caused by our simple data augmentation approach. Nevertheless, the model is able to detect the frauds in the validation set, this is what's important. Figure 1 shows the training loss for the MLP. In the first few iterations the loss is decreasing rapidly, but this trend is slowing down. This is the typical loss for a MLP.

The decision tree has a validation accuracy of 99.96%. The train accuracy is about...

As a naive baseline predictor we present a model that only outputs 0, means no fraud. Because the dataset is imbalanced and data augmentation does not affect the validation set there are only maximum only 0.15% frauds in the validation set. Thus the naive predictor has a accuracy of at least 99.85%. Both implementations are significantly better than the naive baseline predictor. The decision tree and random forest classification are able to detect the frauds.

- Describe the performance of your model (in terms of the metrics for your dataset) on the training and validation sets with the help of plots or/and tables.
- You must provide at least two separate visualizations (plot or tables) of different things, i.e. don't use a table and a bar plot of the same metrics. At least three visualizations are required for the 3 person team.

## IV. DISCUSSION

As mentioned above, our models are significantly better than a model that only predicts no fraudulent transactions.

The decision tree can pretty accurately classify both fraud and non-fraud. In testing it mispredicted 13 non-frauds of 45493 total and 1 frauds of 73 total. This results in a testing accuracy of 0.9996927735960851 which is way above the baseline predictor.

A reason we think the model performed so well is that frauds usually have something that makes them stick out. This enables the decision tree to learn which thresholds for features indicate a fraud. This theory is also supported by the structure of the decision tree as it only uses about a third of the features defined in the dataset. The model is also remarkably small coming in at a little more than 2KB in pure text form.

We also tried removing the timestamp from the dataset since the time of a transaction likely does not hold any information about its fraudulence but it did not improve the model's accuracy and caused problems with the test harness on JupyterHub.

The neural network approach did not work as well as expected. Although we have a high accuracy, the decision tree has better results. Further testing with different numbers of hidden layers in combination with further hyperparameter tuning could lead to better results. Due to limited time we were not able to test this.

- Analyze the results presented in the report (comment on what contributed to the good or bad results). If your method does not work well, try to analyze why this is the case.
- Describe very briefly what you tried but did not keep for your final implementation (e.g. things you tried but that did not work, discarded ideas, etc.).
- How could you try to improve your results? What else would you want to try?

## V. CONCLUSION

For decision trees without pruning the testing accuracy is high per default since the tree can usually perfectly fit the data.

- Finally, describe the test-set performance you achieved. Do not optimize your method based on the test set performance!
- Write a 5-10 line paragraph describing the main take-away of your project.